

# Optimisations and Parallelism of d2q9-bgk.c

George Herbert  
cj19328@bristol.ac.uk

February 20, 2022

## Abstract

d2q9-bgk.c implements the Lattice Boltzmann methods (LBM) to simulate a fluid density on a lattice. This report outlines the techniques I utilised to optimise and parallelise d2q9-bgk.c, as well as a detailed analysis of those techniques. To do so, this report is split into several sections corresponding to different iterations of my code.

## 1 Original code

I compiled the original d2q9-bgk.c using the GNU Compiler Collection (GCC) with the following command:

```
gcc -std=c99 -Wall d2q9-bgk.c -lm -o  
d2q9-bgk.
```

Table 1: Total time of the original code for test cases of different sizes

Test Case Size	Time (s)
$128 \times 128$	0
$128 \times 256$	0
$256 \times 256$	0
$1024 \times 1024$	0

Figure 1 contains the total time to initialise, compute and collate each of the test cases when running the ELF file. It was important to measure the original code, so that I could quantify the performance improvements of my latter implementations. I measured each of the total times by taking an average of 10 runs on BlueCrystal Phase 4's (BC4's) compute nodes. Each of BC4's compute nodes is a Lenovo nx360 M5, which contains two 14-core 2.4 GHz Intel E5-2680 v4 (Broadwell) CPUs and 128 GiB of RAM [1]. I took an average of multiple runs because of the variation between runs, which exists due to the inconsistent performance of compute nodes.

## 2 Serial optimisations

### 2.1 Changes

I changed from GCC to the Intel® C Compiler, since it usually provides better optimised code for Intel processors. Furthermore, I compiled my code with the `Ofast` flag, which set aggressive options to improve the speed of my program, including 03 optimisations and aggressive floating point optimisations [2].

I applied loop fusion to decrease the number of memory accesses in my code. Previously, the entire grid was iterated over in four sequential procedures within each timestep: `propagate`, `rebound`, `collision` and `av_velocity`. By fusing the four loops in these procedures into one, I was able to drastically decrease the number of memory accesses, thereby improving the performance of my program.

Moreover, as a result of implementing the loop fusion, I was also able to implement a pointer swap to eliminate unnecessary swapping between the `cells` and `tmp_cells` arrays. New values of cells are written to the `cells_new` array (renamed from `tmp_cells` to more accurately describe its new purpose). At the end of each timestep, the memory locations the `cells_new` and `cells` variables are pointing to are swapped.

I also improved the arithmetic within each timestep to improve the performance.

### 2.2 Results

Table 2: Total time of the serial optimised code for test cases of different sizes

Test Case Size	Time (s)
$128 \times 128$	0
$128 \times 256$	0
$256 \times 256$	0
$1024 \times 1024$	0

### 3 Vectorization

#### 3.1 Changes

#### 3.2 Results

Table 3: Total time of the vectorized code for test cases of different sizes

Test Case Size	Time (s)
$128 \times 128$	0
$128 \times 256$	0
$256 \times 256$	0
$1024 \times 1024$	0

### 4 Parallelism

#### 4.1 OpenMP

#### 4.2 Results

Table 4: Total time of the parallelised code for test cases of different sizes

Test Case Size	Time (s)
$128 \times 128$	0
$128 \times 256$	0
$256 \times 256$	0
$1024 \times 1024$	0

## References

- [1] *BlueCrystal technical specifications*. URL: <https://www.bristol.ac.uk/acrc/high-performance-computing/hpc-systems-tech-specs/> (visited on Feb. 19, 2022).
- [2] *Alphabetical List of Compiler Options*. June 12, 2021. URL: <https://www.intel.com/content/www/us/en/develop/documentation/cpp-compiler-developer-guide-and-reference/top/compiler-reference/compiler-options/alphabetical-list-of-compiler-options.html> (visited on Feb. 20, 2022).